

Large Language Models – Based Test Case Generation

A Comprehensive Analysis of How AI is Revolutionizing Automated Testing

NIKHIL YATES

```

,; - = + * % # $ % @ * = - - = + * % * + = - ; .
, - = + # # * + = - = + - ; . = , - ; . ; + = - - + # @ # * + - ,
, = # # + - = # - - ; . . . . ; . - % + = ; . . . . ; + - ; . - = # * # = ;
, - * # + ; ; # - ; . . . . + = , ; * ; . . . . = ; . . . . = - ; . . . . ; - % + * * -
= * + = ; . . . . = * ; . . . . # - ; . . . . + = , ; . . . . # * , ; . . . . = - - - ; . . . . - # # -
; % @ = - ; . . . . * = , ; . . . . - + - , + % + - ; . . . . - = # % = - # ; . . . . - = ; . . . . ; % % -
= @ + , ; . . . . - # = - - - ; . . . . ; * @ - , ; . . . . = % + ; . . . . - = - - , ; . . . . # * ; . . . . - = * % * ,
+ ; . . . . + = ; . . . . - = % * ; . . . . ; . . . . ; % , ; . . . . - = , ; . . . . ; % : ; . . . . ; * + ; . . . . - * -
- @ # # - , ; . . . . + , ; . . . . - + % % ; . . . . % + ; . . . . - , ; . . . . + ; . . . . * ; . . . . - - - - - = @ +
# = * ; . . . . ; = * - , ; . . . . - - - - = + + = - + - , % , ; . . . . - * , ; . . . . ; * - , # ; . . . . ; % : , = # - ; + * ,
; % = + ; . . . . ; . . . . + + - ; . . . . - , ; . . . . - = , ; . . . . - * @ % * * * % + , ; . . . . * % ; . . . . = + + - ; % #
, @ # - , ; . . . . - * % * - , ; . . . . ; # ; . . . . + ; = * * = ; . . . . - - ; . . . . = % # = , = , = * @ +
% % ; . . . . - = # - ; . . . . ; + - ; - @ - - - ; . . . . = @ # = ; . . . . ; . . . . + # , = - - ; . . . . - + @ , + , - ;
- , ; . . . . + * = ; . . . . ; = # ; . . . . # + = + @ + - , ; . . . . ; . . . . ; = # + , ; . . . . ; + : ; # , = ; @ *
* % = ; + # ; . . . . ; . . . . - # % + , ; . . . . ; # ; . . . . - + * + = - - ; . . . . ; - * , ; . . . . - + ; . . . . * + + @
+ % + = * * + = - ; . . . . = # # = , ; . . . . - = + , ; . . . . = % = ; . . . . - + = - = - - - - ; . . . . = , ; . . . . * @
, = * % @ * * * * # + ; . . . . - * , ; . . . . ; * - ; . . . . # = - ; . . . . ; . . . . ; . . . . ; . . . . ; * %
; . . . . - = = + @ , ; . . . . - = % - ; . . . . ; * * , ; . . . . - ; . . . . ; . . . . - + = , ; . . . . + * = - - - ; . . . . = -
* * , ; . . . . ; @ , ; . . . . ; . . . . - % + - ; * = - - - - ; . . . . ; . . . . - + , ; . . . . ; - - ; . . . . ; * @ +
+ # - , ; . . . . ; + , ; . . . . - + # - ; . . . . - % = , ; . . . . - * @ % # # * * # % + , ; . . . . ; . . . . + % =
, % @ = ; . . . . ; # - , ; . . . . + , ; . . . . = % % # * % + # @ % # * * * * * # % * * * * + ,
- * * + = - - - - - + * # # @ * # @ - + + + * + % + * * = * + + + + @ * , . . . .
; . . . . - = + + * + + = - ; + + + + = # @ @ # % # # + = = + + + + + # @ * ,
; . . . . - + @ # # + = # * * # * + * # % # * + ;
, + # = - - - # @ ; . . . . ;
, + # - ; + * ,
; # * ; + #
= # * + -

```

Automating Code Generation



AGENDA

- Background and Motivation
 - Code Automation
 - GPT-4
 - Pynguin
- Experiment Design
- Findings
 - The Data
 - Subjective Interpretations
- Discussion and Applications
- Q&A

Current General Exploratory Research

RANDOM TESTS VS COVERAGE-BASED TESTS

- Random approach to test case generation is stronger than the coverage-based approach
- Random approach benefits:
 - Simplicity
 - Minimizes bias
 - Performs well in contexts with large input sizes

EMPIRICAL STUDIES, LLMS

- LLM-based test generation tools are incredibly powerful
 - Tests are realistic
 - Tests could contain assertions
- LLM's are **not** expected to cover edge cases
 - Training data does not encompass non-typical use cases
- Transformer architecture drives high performance on sequential (textual) data
 - LLMs are better for generation regression tests

HUMAN-ORIENTED TEST GENERATION

- All studies complement the 'human' aspect of LLM-based tests
- Reports that include any non-empirical observations do not incorporate them directly into the study
 - Usually an additional statement/comment

LACK OF CRITICISM

- Critical observations about the relationship between empirical and non-empirical qualities of LLM-based tests
- Lack of critical research comparing qualitative attributes of traditional tools and LLM-based tools

Generative Pre-Trained Transformer 4

**175
BILLION Parameters**

Compared to 6 billion in GPT-3.5

- Out-of-box large language model that boasts accuracy, elevated understanding, sophisticated response abilities
- A new animal compared to its predecessors
- New tool = research opportunities
- General Research
 - Strong ability to contextualize input (transformer architecture)
 - Low edge-case inclusion in GPT-3.5
 - GPT-4 stresses the importance of random-case generation
 - Output quality:
 - Clear
 - Relevant
 - Limited to the training data
- Empirical Test Generation Study (*Shäfer et al.*)
 - Realistic
 - No expectation for edge-case coverage
 - Limited qualitative analysis

Pynguin

- Background
 - Unit test generation suite
 - First tool that generates unit tests for dynamically-types languages
 - Currently ***not*** a production-level tool
- Dynamic Generation Approach
 - Generates a basic test
 - Iteratively runs the program against the test and changes the test to improve code coverage
 - Mutants
- Performance limitations increase with codebase complexity
- Most popular *dynamic* test generation tool for python

Experiment Design

Experiment Overview

OBJECTIVE

- Assess the quality of the tests that the LLM-based tool produces
- Compare the LLM-based tool with the traditional tool
 - Performance differences
 - Non-performance differences

DATASET

- TensorFlow GitHub repository
- Rich, diverse data
- Strong and practical codebase

EVALUATION METRICS

- Quantitative *and* Qualitative
- Filling the gaps of existing research
- Emphasis on the relationship between KPIs, not just face value

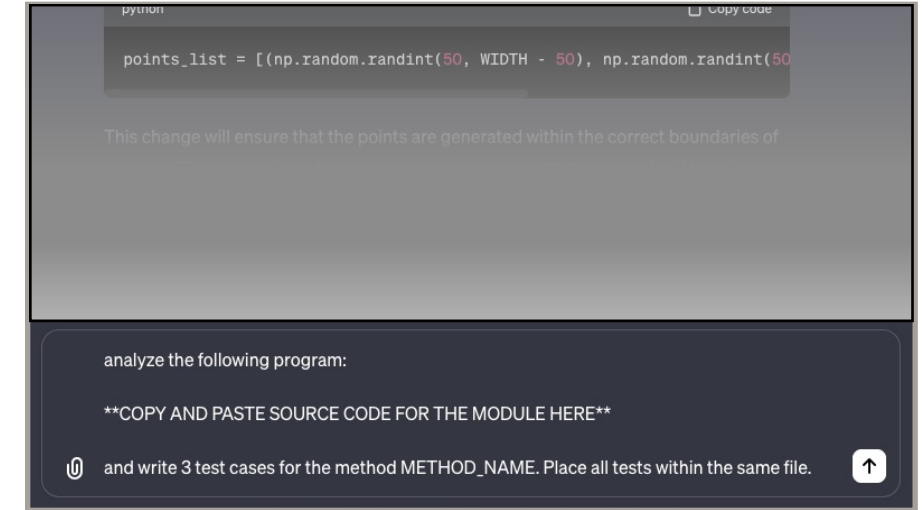
RESEARCH QUESTIONS

- To guide research and experimentation
- The *essence* of the project

Generation Process

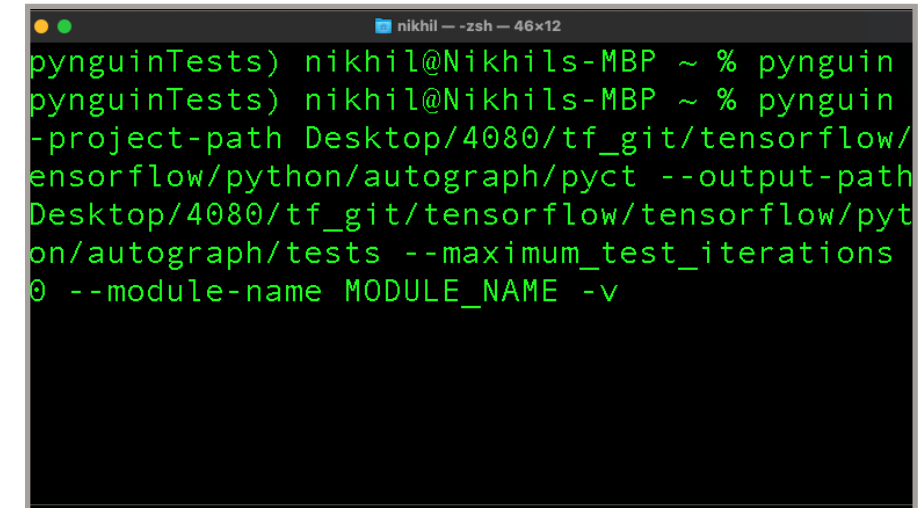
> GPT-4

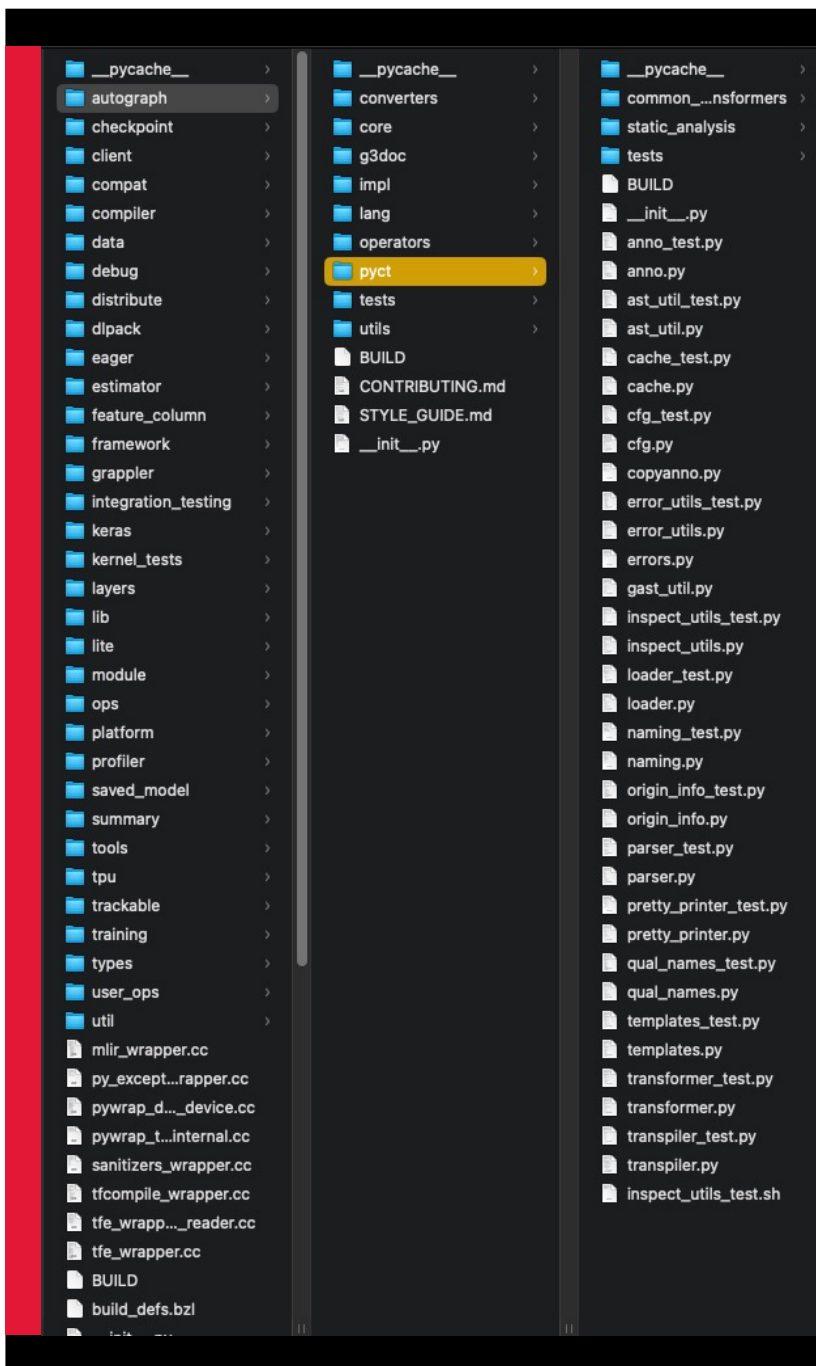
1. Copy the entire python module into the system
2. Method specification
3. *pytest* runnable
4. Repository compatibility*



> Pynguin

1. Configuration analysis
2. Standard execution
3. Repository compatibility*





Dataset information

- TensorFlow/python/autograph/pyct
- 50 methods
- Diversity across the data:
 - methods type (instance, class, static methods)
 - Variable functionality
- Randomness Motivation:
 - Edge case generation
 - Manifested function: LLM training
 - Increased understanding of future code
 - Benchmarking
 - Validation
 - Data richness

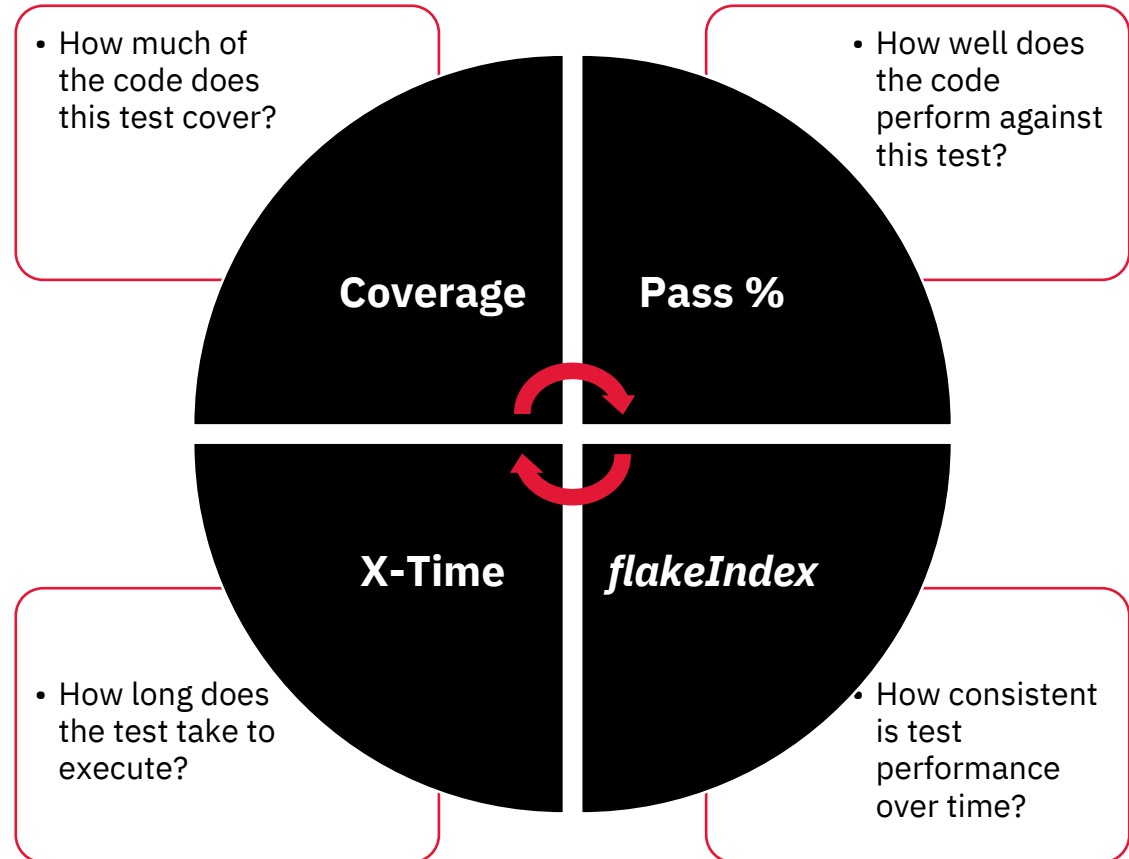
Quantitative Evaluation Metrics

➤ **quantScore**

- Summation of quantitative KPI scores
- Each KPI is judged using a 2-point grading scale
- */8

➤ **flakeIndex**

- Score card to assess the consistency of tests
- Quantitative KPIs are co-dependent
- Specific tests
 - Function of the method



Qualitative Evaluation Metrics

› **qualScore**

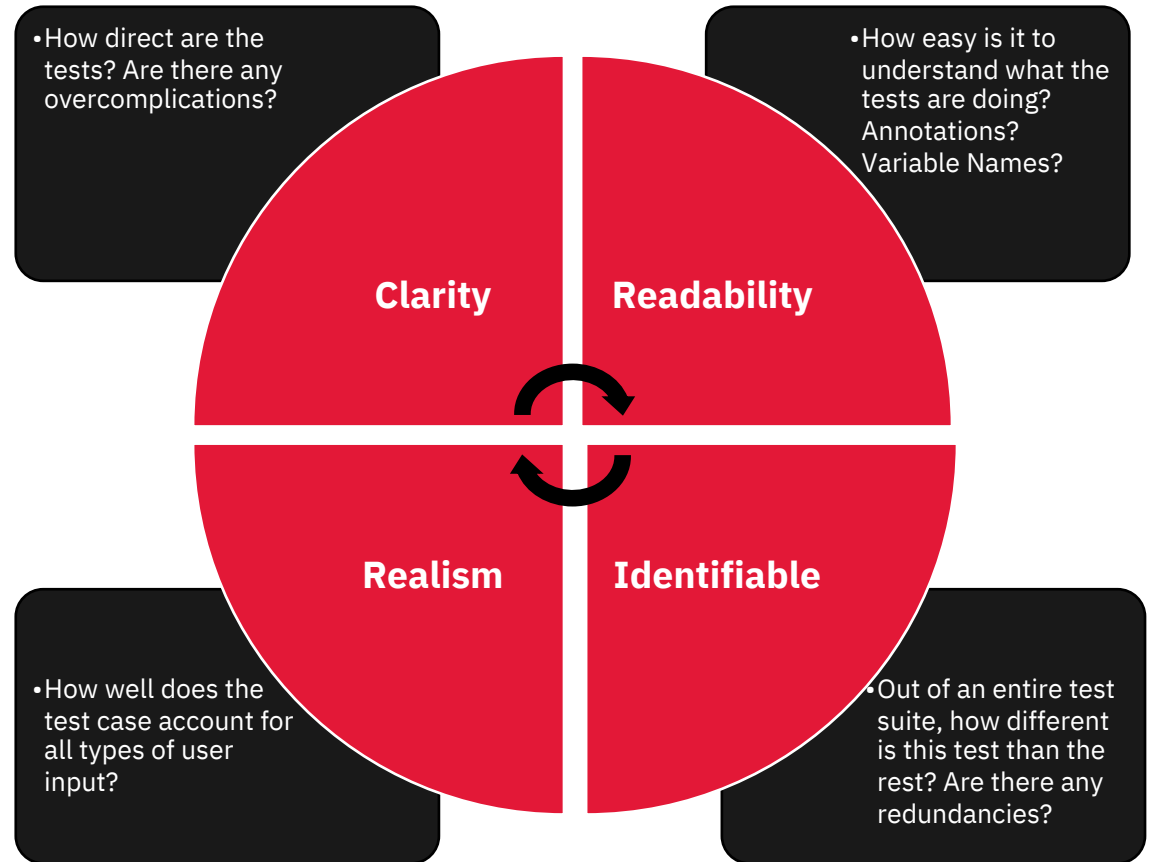
- Summation of qualitative KPI scores
- Each KPI is judged using a 2-point grading scale
- */8

› **Unique nature** of qualitative KPIs

- Subjectiveness
- Independence
- Extensiveness

› Qualitative KPIs are ‘complete’ but inconsistent

- Structure of the tool’s output
 - GPT-4 output can be extensively modified



Guiding Questions

RQ1



(Performance) How well does the LLM-based test generation tool perform with respect to the criteria for strong test cases?

RQ2



(Competence) To what extent does the traditional and LLM-based test generation tools differ quantitatively? Qualitatively?

Result Analysis

Initial Observations

➤ Precision Discrepancies

Tool	Generated Test Files	Generated Test Cases
Pynguin	5	33
GPT-4	50	150

➤ Experiment Integrity

- 1:1 Comparisons

GPT-4 Quantitative Analysis – RQ1

Tool	Coverage	Pass %	Exec. Time (s)	Flakiness
GPT-4	76.4%	70%	6.2	0.0
	1.5/2 <ul style="list-style-type: none">- Complex dataset- Rigidity of	2/2 <ul style="list-style-type: none">- NO xFail tests generated- Testing thoroughness	1.8/2 <ul style="list-style-type: none">- With respect to dataset complexity, complexity of tests, and test coverage	2/2 <ul style="list-style-type: none">- Consistency over 3 rounds of testing

quantScore = 7.3

GPT-4 Qualitative Analysis – RQ1

Tool	Clarity	Readability	Realism	Identifiable
GPT-4	2	1.7	1.9	2
	<ul style="list-style-type: none">▪ Tests were clear and direct▪ The tests did not do unnecessary work (i.e., irrelevant testing)	<ul style="list-style-type: none">▪ Well annotated▪ Appropriate naming conventions▪ Some test files contained minimal redundant comments	<ul style="list-style-type: none">▪ Very accurate in terms of actual use case▪ Some cases did not accommodate arbitrary user input	<ul style="list-style-type: none">▪ No "over-testing"▪ LLM did not generate tests that were similar to previously generated tests▪ All test cases were distinct

qualScore = 7.6


Comparing Tools – RQ2

QUANTITATIVE COMPARISON					
Tool	Coverage	Pass%	Time (s)	Flakiness	<i>quantScore</i>
GPT-4	76.4%	70%	6.2	0.0	7.3
Pynguin	64.2%	21%	5.7	0.1	5.2

QUALITATIVE COMPARISON					
Tool	Clarity	Readability	Realism	Identifiable	<i>qualScore</i>
GPT-4	2	1.7	1.9	2	7.6
Pynguin	0.9	1.1	1.5	1.2	4.7



Pass Frequency



Examples

```
1 import pytest
2 import gast
3 from tensorflow.python.autograph.pyct.ast_util import parallel_walk
4
5 class TestParallelWalk:
6
7     # Test Case 1: Parallel walk with identical ASTs
8     def test_identical_ast(self):
9         node1 = gast.parse("a = 1 + 2")
10        node2 = gast.parse("a = 1 + 2")
11        for n1, n2 in parallel_walk(node1, node2):
12            assert type(n1) == type(n2), "Nodes should be of the same type"
13
14    # Test Case 2: Parallel walk with different ASTs
15    def test_different_ast(self):
16        node1 = gast.parse("a = 1 + 2")
17        node2 = gast.parse("b = 3 + 4")
18        with pytest.raises(ValueError):
19            for _ in parallel_walk(node1, node2):
20                pass
21
22    # Test Case 3: Parallel walk with different structures
23    def test_different_structure(self):
24        node1 = gast.parse("a = 1 + 2")
25        node2 = gast.parse("for i in range(10): a = i")
26        with pytest.raises(ValueError):
27            for _ in parallel_walk(node1, node2):
28                pass
29
30    # If you wish to run the tests directly using this script
31    if __name__ == "__main__":
32        pytest.main()
33
```

```
1 # Test cases automatically generated by PyPenguin (https://www.pyPenguin.eu).
2 # Please check them before you use them.
3 import pytest
4 import tensorflow.python.autograph.pyct.inspect_utils as module_0
5 import inspect as module_1
6 import tokenize as module_2
7 import ast as module_3
8
9
10 def test_case_0():
11     none_type_0 = None
12     var_0 = module_0.isconstructor(none_type_0)
13     assert var_0 is False
14
15 @pytest.mark.xfail(strict=True)
16 def test_case_1():
17     none_type_0 = None
18     var_0 = module_0.islambda(none_type_0)
19     module_0.getqualifiedname(none_type_0, none_type_0, none_type_0)
20
21
22 def test_case_2():
23     none_type_0 = None
24     var_0 = module_1.istraceback(none_type_0)
25     var_1 = module_0.isnamedtuple(var_0)
26     assert var_1 is False
27
28
29 @pytest.mark.xfail(strict=True)
30 def test_case_3():
31     bool_0 = True
32     var_0 = module_0.isbuiltin(bool_0)
33     assert var_0 is True
34     var_1 = module_2.group()
35     module_0.getnamespace(var_1)
36
37
38 @pytest.mark.xfail(strict=True)
39 def test_case_4():
40     none_type_0 = None
41     var_0 = module_0.getfutureimports(none_type_0)
42     var_1 = module_0.isbuiltin(var_0)
43     assert var_1 is False
44     var_2 = module_0.isnamedtuple(none_type_0)
45     assert var_2 is False
46     module_0.getqualifiedname(none_type_0, none_type_0, none_type_0, none_type_0)
47
48
```

Conclusion

RQ1



(Performance) How well does the LLM-based test generation tool perform with respect to the criteria for strong test cases?

RQ2



(Competence) To what extent does the traditional and LLM-based test generation tools differ quantitatively? Qualitatively?

Applications and Discussion

Further Research and Applications

- Project focused on developing test cases for the python code base:
 - What about a statically typed language like Java?
- On 50 test cases from the tensorflow/java/ directory:
- Strong *quantScore* of 7.1 and *qualScore* of 7.5
- Comparison against a traditional unit-test generation tool like Randoop, Palus?

- The benefits of test-driven code development
 - Robustness
 - Comprehension
 - Validity
 - Transformation of code
 - Collaboratory benefits
 - And more!
- Test-drive development in schools
 - A meaningful piece of criteria

Questions

Thank You!

Acknowledgements to Prof. Song Wang